

Linear Operators in a Generative Adversarial Network

Jasper Hill

February 23, 2020

1 Introduction

The quest to emulate human-like intelligence in computers has been aided significantly by the employment of convolutional neural networks (CNNs). In accordance with its name, a CNN features one or more layers that act to convolve each input channel of a tensor with a kernel to generate some number of output channels with altered shapes. Though the success of CNNs is unambiguous, their mathematical foundation seems a bit unsound. The purpose of this experiment is to develop an alternative set of tools and to investigate their efficacy within the context of a generative adversarial network.

2 Convolutional Layers

A specific element of the convoluted tensor, A' , can be expressed in terms of the input tensor, A , and convolution operator, \hat{C} :

$$A'_{ij}{}^{(n_o)} = (\hat{C}A)_{ij}{}^{(n_o)} = \left(\sum_{n_i} \hat{C}^{(n_o)} * A^{(n_i)} \right)_{ij} = \sum_{n_i} \sum_{kl} c_{kl}^{(n_o)} A_{i-kj-l}^{(n_i)}. \quad (1)$$

Here, n_o and n_i are the output and input channel indices respectively, and $*$ is the Hadamard operator in $K \times L$ space. It is clear from the rightmost expression that each transformed tensor element encodes only a finite amount of information from the original tensor. In order to correlate tensor elements respectively separated by heights or widths greater than K or L , multiple convolution layers must be successively applied. For this reason, attaining highly-sophisticated functionality, such as image segmentation or style transfer, requires increasingly many layers to couple each of the original tensor elements with one another.

3 Linear Operators

Unlike the convolutional layer, which employs the Hadamard product between its kernels and the input matrices, the linear operators utilized in this experiment facilitate a transformation on the full space of the input tensors:

$$\begin{aligned}
A'_{ij}^{(n_o)} &= (\hat{O}A\hat{P})_{ij}^{(n_o)} \\
&= \left(\sum_{n_i}^{N_i} \hat{O}^{(n_o, n_i)} A^{(n_i)} \hat{P}^{(n_o, n_i)} \right)_{ij} \\
&= \sum_{n_i}^{N_i} \sum_{kl}^{MN} O_{ik}^{(n_o, n_i)} A_{kl}^{(n_i)} P_{lj}^{(n_o, n_i)}
\end{aligned} \tag{2}$$

for $A \in \mathbb{R}^{M \times N}$ with $O \in \mathbb{R}^{M' \times M}$ and $P \in \mathbb{R}^{N \times N'}$. The output, A' , then, will contain N_o matrices of dimension $M' \times N'$. Importantly, each element of the transformed tensor now encodes all information from the input. In theory, all necessary spatial correlations can be learned with only one layer.

4 Tensor Flattening

When a network is used to generate scalar information from two-dimensional images, the common strategy is to employ a series of convolutional layers, and, when a sufficient number of output channels has been produced, the resulting tensors are unpacked so that they are of rank 1 (i.e., vectors). It is assumed then, that any important spatial correlation among these vector elements can be learned by the input weights of the remaining dense layers. A more mathematically-rigorous way to ensure this is to integrate the image over vectors \vec{u} and \vec{v} :

$$\vec{y} = \vec{u}^\dagger A \vec{v} \tag{3}$$

This operation converts the rank-3 tensor, A , into the vector, \vec{y} , whose elements are given by

$$y^{(n_i)} = \sum_{ij}^{MN} u_i^{(n_i)} A_{ij}^{(n_i)} v_j^{(n_i)}. \tag{4}$$

In this way, the 2-dimensional structure of the input matrices is integrated into each element of the flattened tensor.